



OpenEnv Hackathon



Northflank Docs | OpenEnv Hackathon

Northflank is the developer platform you'll use to deploy your work, using compute provided by Coreweave.


tl;dr

- To get started, you first need to request access to the compute by filling out [this form](#).
- Every team gets an H100.
- You will receive an email from Northflank from services@northflank.com.
- Accept your Northflank Team invite in your email, this will open your Northflank Team dashboard


- In your Northflank Project called **hackathon**, you can create a Northflank Service (container) with access to your H100. You can follow our guides to deploy :
 - **Recommended:** a Jupyter notebook w/ PyTorch preinstalled – 📄 [Jupyter Notebook with PyTorch](#)
 - A generic Ubuntu image – 📄 [Shell access \(VM-like experience\)](#)
 - Your existing GitHub repository using Docker – 📄 [CI/CD Python application](#)
 - Or any Public Docker image – 📄 [Common AI Docker Images](#)
- You can then follow our guide for 📄 [Using SSH to access your Services](#) to connect to your instance w/ VSCode/Cursor/etc
- In case you need to troubleshoot something, reference the 📄 [Use and troubleshoot your deployments](#) or mention anyone from the Northflank team in the OpenEnv Discord channel

In-depth guide

1. Your team will have a Northflank team and project associated to you. You'll receive an invite for your team to your email address:



- Once you've made an account and accepted the invite, you'll be presented with the Northflank home page. Here you can find your project called **Hackathon**. You can click on it and this will navigate you to the project dashboard.




- Your team will have access to a GPU node with a single NVIDIA H100 for your task.
- You can then deploy your code to Northflank, for example
 - Recommended:** a Jupyter notebook w/ PyTorch preinstalled – [🔗 Jupyter Notebook with PyTorch](#)
 - A generic Ubuntu image – [🔗 Shell access \(VM-like experience\)](#)
 - Your existing GitHub repository using Docker – [🔗 CI/CD Python application](#)
 - Or any Public Docker image – [🔗 Common AI Docker Images](#)



You can only have 1 GPU service, if you need more, talk with the organizers to see if there's more capacity available. You can scale down your existing service if you want to try something else.

Extra information

- The development sandboxes are running `Ubuntu 22.04.5` and are based off the `pytorch/pytorch:2.8.0-cuda12.6-cudnn9-runtime` Docker image. Any system packages you install or changes you make to your environment won't be persisted upon reboot.
- You can install the Northflank CLI on your machine (<https://northflank.com/docs/v1/api/use-the-cli>) and easily upload/download files to/from the sandbox, e.g. to visualise ncu/nsys reports using the NVIDIA GUI tools
 - You can install the CLI on your laptop using `npm` or `yarn` - `yarn global add @northflank/cli` `npm i -g @northflank/cli`
 - After you've installed it, you can run `northflank login` and follow the instructions to authenticate your CLI with your Northflank team.
- You can also view the container logs and metrics of the shell, in the Northflank UI, by clicking on the respective buttons on the Service's page:



7 things to know when deploying AI workloads



Select enough **CPU and memory** to run your workload - this is the **compute plan** in the resources section when you create a new service. You can also **select a GPU** and the number of GPUs to assign to that service. Your project has access to **128 vCPU, 2 TB memory, and 1xH100 GPU**

[Scale CPU and memory](#)

 Increase the **ephemeral storage** for your builds and deployments. You will likely experience container evictions otherwise. 5-10GB is usually sufficient, but it depends on the size of your model. This is under **advanced resource options** in the **resources section** when you create a new service. You might also need to increase the SHM size for some applications.

 [Scale ephemeral storage](#)

 If you want to deploy an image but it keeps restarting you'll need to set the runtime command in the **advanced section** when creating a service, or the **CMD override** view in a service. Select **custom entrypoint and command** with the entrypoint as `/bin/bash -c` or `/bin/sh -c` and the command as `sleep 3d`. You can also use this to set a custom command to run a script or application when your container starts.

 [Override command or entrypoint](#)

 **Containers are ephemeral**—all changes made inside them are erased when they restart. To permanently save your models or data, you'll need to **add persistent storage**. The mount path corresponds to the container's file system.

 When you've created a build or deployment service you can select a container from the **observability dashboard** to view **logs, metrics, and health** information. This can help you debug issues.

 [Monitor containers](#)

 You can **click the terminal button** on a running container to gain shell access, or click shell in the container view. This will start a prompt with the default user for the image. The image will only have access to programs that it was built with, for example some images may not have a package manager or text editor. You can also install the Northflank CLI to access a container's shell on your machine.



You can expose your application to the internet in the **networking section** when creating a deployment, or from a service's **ports and DNS** page. You can use this to access web interfaces or APIs served by your deployments. Only HTTP can be publicly exposed, and ports must correspond to ports exposed by your application. Alternatively, you can use the Northflank CLI to securely forward your application to your local machine.