Northflank Docs | OpenEnv Hackathon

Northflank is the developer platform used to deploy hackathon workloads using compute provided by CoreWeave. This document summarizes the onboarding and deployment steps for teams participating in the OpenEnv hackathon.

Quick Start (tl;dr)

- Request compute access by filling out the provided Google Form.
- Each team receives access to one NVIDIA H100 GPU.
- You will receive an invite email from services@northflank.com.
- Accept the invite to access your Northflank team dashboard.
- In the project named 'hackathon', create a service (container) with GPU access.

Deployment Options

- Jupyter Notebook with PyTorch preinstalled.
- Generic Ubuntu image with shell access (VM■like experience).
- Deploy an existing GitHub repository using Docker and CI/CD.
- Deploy any public Docker image.

Extra Environment Information

- Development sandboxes run Ubuntu 22.04.5.
- They are based on the Docker image pytorch/pytorch:2.8.0-cuda12.6-cudnn9-runtime.
- Environment changes inside containers are not persisted after restart.
- You can install the Northflank CLI to upload/download files and access services.
- Authenticate the CLI with 'northflank login'.

Important Deployment Notes

- Choose sufficient CPU and memory for workloads. Projects typically have up to 128 vCPU, 2 TB RAM, and 1 H100 GPU.
- Increase ephemeral storage (5–10GB recommended) to avoid container eviction.
- If containers repeatedly restart, set a runtime command such as '/bin/bash -c sleep 3d'.
- Containers are ephemeral—use persistent storage to keep models or datasets.
- Use logs and metrics from the observability dashboard to debug deployments.

- You can open a shell directly into running containers.
- Expose applications via HTTP ports or use the CLI to securely forward services to your local machine.